

Threads

Tran, Van Hoai

Faculty of Computer Science & Engineering
HCMC University of Technology

E-mail: hoai@hcmut.edu.vn
(*partly based on slides of Le Thanh Van*)

1 Thread and its benefits

2 Multithreading models

1 Thread and its benefits

2 Multithreading models

- An application normally has several controls


A word-processor has a control on mouse input, a control for keyboard, a control for function completion,

...


- The model “an application = a process” does not catch up with multiprocessor environment

A modern processor has multiple cores

Single vs. Multithreaded processes




single-threaded process




multithreaded process

Single vs. Multithreaded processes




single-threaded process



multithreaded process

With the above model, **one** server (by one process) can service **several concurrent** requests.



Benefits of multithreading model

- **Responsiveness:** a program (process) **continues running** even if a part of it is blocked or is performing a lengthy operation.
- **Resource sharing:** By defaults, threads share memory and resources of its process \Rightarrow **same address-space**.
- **Economy:** Resource allocation, context-switching are **time-consuming**. Threads do it **more economically**.
- **Scalability:** threads may be running **in parallel** on different processing cores.

Benefits of multithreading model

- **Responsiveness:** a program (process) **continues running** even if a part of it is blocked or is performing a lengthy operation.
- **Resource sharing:** By defaults, threads share memory and resources of its process \Rightarrow **same address-space**.
- **Economy:** Resource allocation, context-switching are **time-consuming**. Threads do it **more economically**.
- **Scalability:** threads may be running **in parallel** on different processing cores.

concurrency vs. parallelism

- Concurrency = many tasks are allowed to make progress
- Parallelism = many tasks can be performed simultaneously

Programming challenges

- **Identifying tasks:** how to divide an application into tasks ?
- **Balance:** how tasks do the same amount of workload ?
- **Data splitting:** how data of tasks to be splitted ?
- **Data dependency:** data surely does not live alone, how they are synchronized ?
- **Testing and debugging:** how to follow many different execution paths ?


Programming challenges

- **Identifying tasks:** how to divide an application into tasks ?
- **Balance:** how tasks do the same amount of workload ?
- **Data splitting:** how data of tasks to be splitted ?
- **Data dependency:** data surely does not live alone, how they are synchronized ?
- **Testing and debugging:** how to follow many different execution paths ?

Textbook

“Many computer science educators believe that [software development](#) must be taught with increased emphasis on [parallel programming](#).”

Do not expect much on multiple threading



Amdahl's Law

$$speedup \leq \frac{1}{S + \frac{1-S}{N}}$$

in which, S : serial portion; N : number of cores (threads).

1 Thread and its benefits

2 Multithreading models

User vs. Kernel threads

User threads

- Thread management done by user-level thread library
- Examples: POSIX Pthreads, Mach C-threads, Solaris threads

Kernel threads

- Thread management done at kernel-level by OS
- Examples: Windows, Linux, Max OS X, Solaris

User vs. Kernel threads

User threads


- Thread management done by user-level thread library
- Examples: POSIX Pthreads, Mach C-threads, Solaris threads

Kernel threads

- Thread management done at kernel-level by OS
- Examples: Windows, Linux, Max OS X, Solaris


A relationship must exist between user threads and kernel threads

Many-to-one




- Mapping **many** user-level threads to **one** kernel thread
- Issues:

Many-to-one




- Mapping **many** user-level threads to **one** kernel thread
- Issues:
 - if a thread is blocked, the **entire** process is blocked too.
 - Unable to run in parallel on multicore systems

One-to-one




- Mapping **each** user thread to **a** kernel thread
- Issues:

One-to-one




- Mapping **each** user thread to **a** kernel thread
- Issues:
 - Creating a user thread means creating a kernel thread \Rightarrow **overhead**.
 - Number of threads is **restricted**
 - **Linux, Windows**

Many-to-many



- Multiplexing many user-level threads to a smaller or equal number of kernel threads
- Issues:

Many-to-many



- Multiplexing many user-level threads to a smaller or equal number of kernel threads
- Issues:
 - Not so many OS implementations apply this model, (Solaris supports)