

Process synchronization

Tran, Van Hoai

Faculty of Computer Science & Engineering
HCMC University of Technology

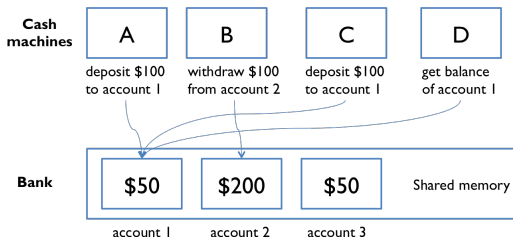
E-mail: hoai@hcmut.edu.vn
(*partly based on slides of Le Thanh Van*)

- 1 Basic concepts
- 2 Critical-section (CS) problem
- 3 Synchronization hardware
- 4 Software tools for synchronization
 - Mutex lock
 - Semaphores
- 5 Synchronization problems

- 1 Basic concepts
- 2 Critical-section (CS) problem
- 3 Synchronization hardware
- 4 Software tools for synchronization
 - Mutex lock
 - Semaphores
- 5 Synchronization problems

Why do we study synchronization ? (1)

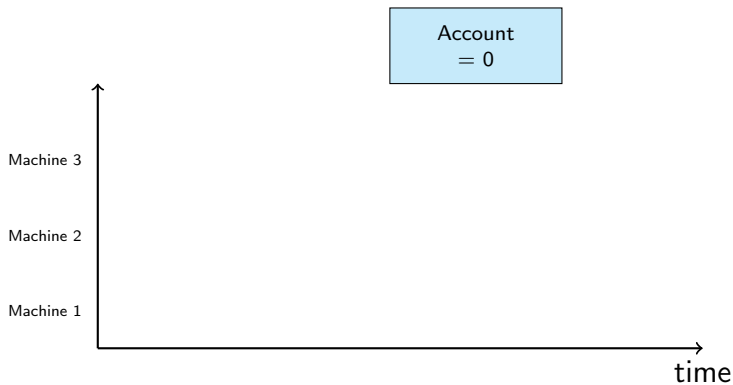
- Different CPU schedulers produce different timings for scheduled processes
- The correctness of a concurrent program **should not depend** on accidents of timing



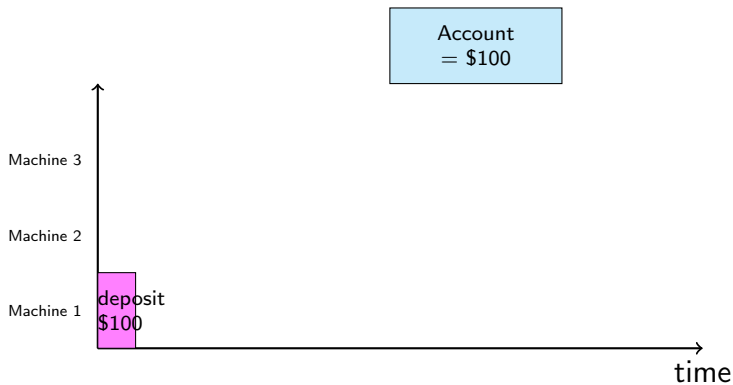
(Source: mit.edu)

Cooperating (concurrent) processes which possibly share a logical address can create the **inconsistency** in shared data

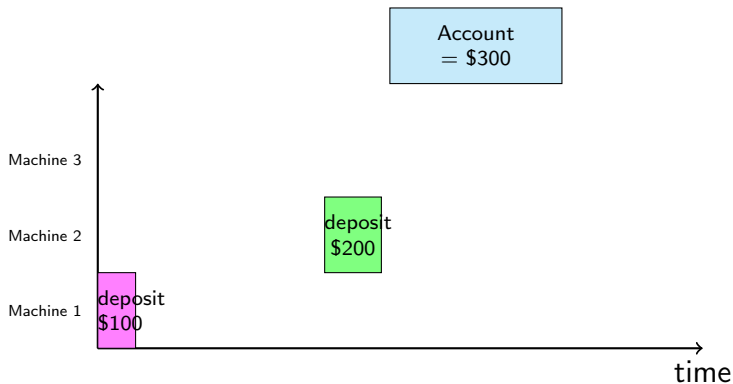
Why do we study synchronization ? (2)



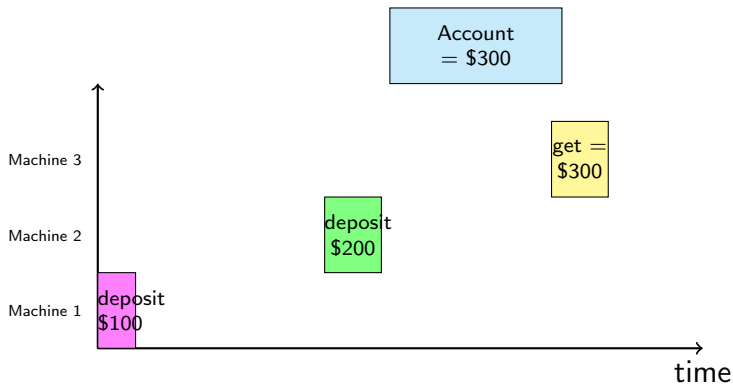
Why do we study synchronization ? (2)



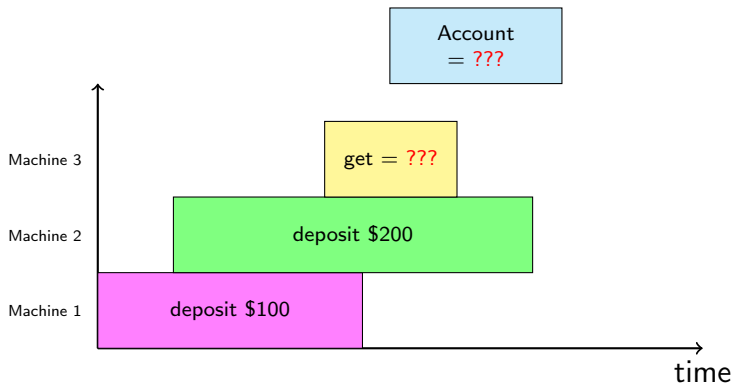
Why do we study synchronization ? (2)



Why do we study synchronization ? (2)



Why do we study synchronization ? (3)



Inconsistency in shared-memory communication (1)

Use a counter to remedy (chữa bệnh) for buffer size of
`BUFFER_SIZE - 1`

Producer

```
item next_produced;
while (true){
    while ( counter == BUFFER_SIZE )
        ; /* do nothing */
    buffer[ in ] = next_produced;
    in = ( in + 1 ) % BUFFER_SIZE;
    counter ++;
}
```

Consumer

```
item next_consumed;
while (true){
    while ( counter == 0 )
        ; /* do nothing */
    next_consumed = buffer[ out ];
    out = ( out + 1 ) % BUFFER_SIZE;
    counter --;
}
```

- The codes seem correct, but not
- Suppose currently counter = 5, then 2 statements counter++ and counter-- executed concurrently
- After those executed, value of counter could be 4, 5, or 6 (inconsistently).
(the expected is 5)

Inconsistency in shared-memory communication (2)

counter++

```
register1 = counter  
register1 = register1 + 1  
counter = register1
```

counter--

```
register2 = counter  
register2 = register2 - 1  
counter = register2
```

Concurrent execution of 2 processes is interleaved in a **low-level sequential** execution in **any arbitrary** order, perhaps as follows.

T_0 :	producer	execute	$register_1 = counter$	($register_1 = 5$)
T_1 :	producer	execute	$register_1 = register_1 + 1$	($register_1 = 6$)
T_2 :	consumer	execute	$register_2 = counter$	($register_2 = 5$)
T_3 :	consumer	execute	$register_2 = register_2 - 1$	($register_2 = 4$)
T_4 :	producer	execute	$counter = register_1$	($counter = 6$)
T_5 :	consumer	execute	$counter = register_2$	($counter = 4$)

Inconsistency in shared-memory communication (2)

counter++

```
register1 = counter  
register1 = register1 + 1  
counter = register1
```

counter--

```
register2 = counter  
register2 = register2 - 1  
counter = register2
```

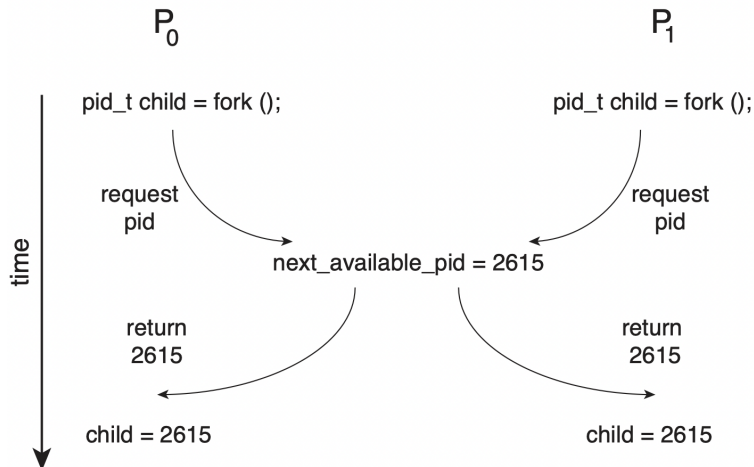
Concurrent execution of 2 processes is interleaved in a **low-level sequential** execution in **any arbitrary** order, perhaps as follows.

T_0 : producer execute $register_1 = counter$ ($register_1 = 5$)

Race condition

Outcome of executing of multiple cooperating concurrent processes depends on **particular order** of their access events

Race condition example: fork()



- 1 Basic concepts
- 2 Critical-section (CS) problem
- 3 Synchronization hardware
- 4 Software tools for synchronization
 - Mutex lock
 - Semaphores
- 5 Synchronization problems

Critical section

Multiple processes should have the following structure to avoid race condition

```
do {  
    entry section  
    critical section  
    exit section  
    remainder section  
} while (true);
```



Critical section

Multiple processes should have the following structure to avoid race condition

```
do {  
    entry section  
        critical section  
    exit section  
        remainder section  
} while (true);
```



Critical section in kernel-mode process

- Kernel-mode processes share some kernel data structures
list of opened files
- 2 approaches to handle critical sections in OS
 - **Nonpreemptive kernel**: kernel-mode process cannot be preempted
→ No race conditions
 - **Preemptive kernel**: kernel-mode process can be preempted
→ more responsive, but prone to critical section

Conditions to critical section

- **Mutual exclusion**: at most one process allowed in critical section
- **Progress**: if no process in critical section and some (one or more) wish to enter critical section, the selection of which (of these processes) will enter cannot be postponed indefinitely
- **Bounded waiting**: a limit on number of times which other processes allowed to enter critical section after a process makes a request and before that request granted.

Algorithm 1 for CS

(assumed for case of 2 processes)

- Processes **do not know** which is in CS
- \rightarrow use a **shared variable** “turn” to **know** which is in CS. (Initially $\text{turn}=0$)
 $\text{turn} = i \Rightarrow P_i$ is in its CS

Process P_i

```
do {  
    while ( turn != i ) ;  
        /* critical section code is here */  
    turn = j;  
        /* remainder section code is here */  
} while(true);
```

Algorithm 1 for CS

(assumed for case of 2 processes)

- Processes **do not know** which is in CS
- \rightarrow use a **shared variable** “turn” to **know** which is in CS. (Initially turn=0)
turn = i $\Rightarrow P_i$ is in its CS

Process P_i

```
do {  
    while ( turn != i ) ;  
        /* critical section code is here */  
    turn = j;  
        /* remainder section code is here */  
} while(true);
```

Satisfies **mutual exclusion**, **bounded waiting**

Algorithm 1 for CS


(assumed for case of 2 processes)

- Processes **do not know** which is in CS
- \rightarrow use a **shared variable** “turn” to **know** which is in CS. (Initially $\text{turn}=0$)
 $\text{turn} = i \Rightarrow P_i$ is in its CS

Process P_i

```
do {  
    while ( turn != i ) ;  
        /* critical section code is here */  
    turn = j;  
        /* remainder section code is here */  
} while(true);
```

Satisfies **mutual exclusion**, **bounded waiting**



Algorithm 1 for CS


(assumed for case of 2 processes)

- Processes **do not know** which is in CS
- use a **shared variable** “turn” to **know** which is in CS. (Initially $\text{turn}=0$)
 $\text{turn} = i \Rightarrow P_i$ is in its CS

Process P_i

```
do {  
    while ( turn != i ) ;  
        /* critical section code is here */  
    turn = j;  
        /* remainder section code is here */  
} while(true);
```

Satisfies **mutual exclusion**, **bounded waiting**



Algorithm 1 for CS


(assumed for case of 2 processes)

- Processes **do not know** which is in CS
- \rightarrow use a **shared variable** “turn” to **know** which is in CS. (Initially $\text{turn}=0$)
 $\text{turn} = i \Rightarrow P_i$ is in its CS

Process P_i

```
do {  
    while ( turn != i ) ;  
        /* critical section code is here */  
    turn = j;  
        /* remainder section code is here */  
} while(true);
```

Satisfies **mutual exclusion**, **bounded waiting**



Algorithm 1 for CS


(assumed for case of 2 processes)

- Processes **do not know** which is in CS
 - use a **shared variable** “turn” to **know** which is in CS. (Initially turn=0)
- turn = i \Rightarrow P_i is in its CS

Process P_i

```
do {  
    while ( turn != i ) ;  
        /* critical section code is here */  
    turn = j;  
        /* remainder section code is here */  
} while(true);
```

Satisfies **mutual exclusion**, **bounded waiting**



Algorithm 1 for CS


(assumed for case of 2 processes)

- Processes **do not know** which is in CS
- use a **shared variable** “turn” to **know** which is in CS. (Initially turn=0)
turn = i \Rightarrow P_i is in its CS

Process P_i

```
do {  
    while ( turn != i ) ;  
        /* critical section code is here */  
    turn = j;  
        /* remainder section code is here */  
} while(true);
```

Satisfies **mutual exclusion**, **bounded waiting**



Algorithm 1 for CS


(assumed for case of 2 processes)

- Processes **do not know** which is in CS
 - use a **shared variable** “turn” to **know** which is in CS. (Initially turn=0)
- turn = i \Rightarrow P_i is in its CS

Process P_i

```
do {  
    while ( turn != i ) ;  
    /* critical section code is here */  
    turn = j;  
    /* remainder section code is here */  
} while(true);
```

Satisfies **mutual exclusion**, **bounded waiting**



Algorithm 1 for CS



(assumed for case of 2 processes)

- Processes **do not know** which is in CS
 - use a **shared variable** “turn” to **know** which is in CS. (Initially turn=0)
- turn = i \Rightarrow P_i is in its CS

Process P_i

```
do {  
    while ( turn != i ) ;  
        /* critical section code is here */  
    turn = i;  
        /* remainder section code is here */  
} while(true);
```

Satisfies **mutual exclusion**, **bounded waiting** , but **not progress**



Algorithm 1 for CS



(assumed for case of 2 processes)

- Processes **do not know** which is in CS
 - use a **shared variable** “turn” to **know** which is in CS. (Initially turn=0)
- turn = i \Rightarrow P_i is in its CS

Process P_i

```
do {  
    while ( turn != i ) ;  
    /* critical section code is here */  
    turn = j;  
    /* remainder section code is here */  
} while(true);
```

Satisfies **mutual exclusion**, **bounded waiting** , but **not progress**



Algorithm 1 for CS



(assumed for case of 2 processes)

- Processes **do not know** which is in CS
- use a **shared variable** “turn” to **know** which is in CS. (Initially turn=0)
- turn = i \Rightarrow P_i is in its CS

Process P_i

```
do {  
    while ( turn != i ) ;  
    /* critical section code is here */  
    turn = j;  
    /* remainder section code is here */  
} while(true);
```

Satisfies **mutual exclusion**, **bounded waiting** , but **not progress**



Algorithm 1 for CS



(assumed for case of 2 processes)

- Processes **do not know** which is in CS
- use a **shared variable** “turn” to **know** which is in CS. (Initially turn=0)
- turn = i \Rightarrow P_i is in its CS

Process P_i

```
do {  
    while ( turn != i ) ;  
    /* critical section code is here */  
    turn = j;  
    /* remainder section code is here */  
} while(true);
```

Satisfies **mutual exclusion**, **bounded waiting**, but **not progress**



Algorithm 1 for CS



(assumed for case of 2 processes)

- Processes **do not know** which is in CS
- use a **shared variable** “turn” to **know** which is in CS. (Initially turn=0)
- turn = i \Rightarrow P_i is in its CS

Process P_i

```
do {  
    while ( turn != i ) ;  
    /* critical section code is here */  
    turn = j;  
    /* remainder section code is here */  
} while(true);
```

Satisfies **mutual exclusion**, **bounded waiting** , but **not progress**



Algorithm 2 for CS

- turn cannot keep the state of readiness into CS
- \rightarrow use a shared variable “flag” to keep the state of all processes
 $\text{flag}[i] = \text{true} \Rightarrow P_i$ is in CS

Process P_i

```
do {  
    flag[i] = true;  
    while ( flag[j] ) ;  
        /* critical section code is here */  
    flag[i] = false;  
        /* remainder section code is here */  
} while(true);
```

Algorithm 2 for CS

- turn cannot keep the state of readiness into CS
- \rightarrow use a shared variable “flag” to keep the state of all processes
 $\text{flag}[i] = \text{true} \Rightarrow P_i$ is in CS

Process P_i

```
do {  
    flag[i] = true;  
    while ( flag[j] ) ;  
        /* critical section code is here */  
    flag[i] = false;  
        /* remainder section code is here */  
} while(true);
```

mutual exclusion, but not bounded waiting, and not progress


Algorithm 2 for CS

- turn cannot keep the state of readiness into CS
- → use a shared variable “flag” to keep the state of all processes
 $\text{flag}[i] = \text{true} \Rightarrow P_i$ is in CS

Process P_i

```
do {  
    flag[i] = true;  
    while ( flag[j] ) ;  
        /* critical section code is here */  
    flag[i] = false;  
        /* remainder section code is here */  
} while(true);
```

mutual exclusion, but not bounded waiting, and not progress




Algorithm 2 for CS

- turn cannot keep the state of readiness into CS
- \rightarrow use a shared variable “flag” to keep the state of all processes
 $\text{flag}[i] = \text{true} \Rightarrow P_i$ is in CS

Process P_i

```
do {  
    flag[i] = true;  
    while ( flag[j] ) ;  
        /* critical section code is here */  
    flag[i] = false;  
        /* remainder section code is here */  
} while(true);
```

mutual exclusion, but not bounded waiting, and not progress




Algorithm 2 for CS

- turn cannot keep the state of readiness into CS
- → use a shared variable “flag” to keep the state of all processes
 $\text{flag}[i] = \text{true} \Rightarrow P_i$ is in CS

Process P_i

```
do {  
    flag[i] = true;  
    while ( flag[j] ) ;  
        /* critical section code is here */  
    flag[i] = false;  
        /* remainder section code is here */  
} while(true);
```

mutual exclusion, but not bounded waiting, and not progress




Algorithm 2 for CS

- turn cannot keep the state of readiness into CS
- → use a shared variable “flag” to keep the state of all processes
 $\text{flag}[i] = \text{true} \Rightarrow P_i$ is in CS

Process P_i

```
do {  
    flag[i] = true;  
    while ( flag[j] ) ;  
        /* critical section code is here */  
    flag[i] = false;  
        /* remainder section code is here */  
} while(true);
```

mutual exclusion, but not bounded waiting, and not progress



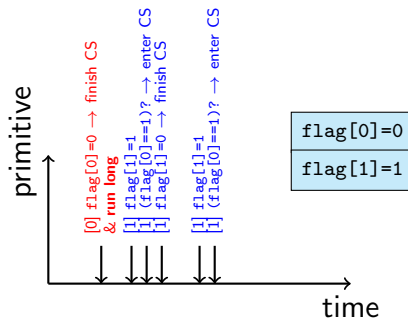
Algorithm 2 for CS

- turn cannot keep the state of readiness into CS
- → use a shared variable “flag” to keep the state of all processes
 $\text{flag}[i] = \text{true} \Rightarrow P_i$ is in CS

Process P_i

```
do {  
    flag[i] = true;  
    while ( flag[j] ) ;  
    /* critical section code is here */  
    flag[i] = false;  
    /* remainder section code is here */  
} while(true);
```

mutual exclusion, but not bounded waiting, and not progress





Algorithm 2 for CS

- turn cannot keep the state of readiness into CS
- → use a shared variable “flag” to keep the state of all processes
 $\text{flag}[i] = \text{true} \Rightarrow P_i$ is in CS

Process P_i

```
do {  
    flag[i] = true;  
    while ( flag[j] ) ;  
        /* critical section code is here */  
    flag[i] = false;  
        /* remainder section code is here */  
} while(true);
```

mutual exclusion, but not bounded waiting, and not progress





Algorithm 2 for CS

- turn cannot keep the state of readiness into CS
- → use a shared variable “flag” to keep the state of all processes
 $\text{flag}[i] = \text{true} \Rightarrow P_i$ is in CS

Process P_i

```
do {  
    flag[i] = true;  
    while ( flag[j] ) ;  
    /* critical section code is here */  
    flag[i] = false;  
    /* remainder section code is here */  
} while(true);
```

mutual exclusion, but not bounded waiting, and not progress





Algorithm 2 for CS

- turn cannot keep the state of readiness into CS
- → use a shared variable “flag” to keep the state of all processes
 $\text{flag}[i] = \text{true} \Rightarrow P_i$ is in CS

Process P_i

```
do {  
    flag[i] = true;  
    while ( flag[j] ) ;  
    /* critical section code is here */  
    flag[i] = false;  
    /* remainder section code is here */  
} while(true);
```

mutual exclusion, but not bounded waiting, and not progress





Algorithm 2 for CS

- turn cannot keep the state of readiness into CS
- → use a shared variable “flag” to keep the state of all processes
 $\text{flag}[i] = \text{true} \Rightarrow P_i$ is in CS

Process P_i

```
do {  
    flag[i] = true;  
    while ( flag[j] ) ;  
    /* critical section code is here */  
    flag[i] = false;  
    /* remainder section code is here */  
} while(true);
```

mutual exclusion, but not bounded waiting, and not progress



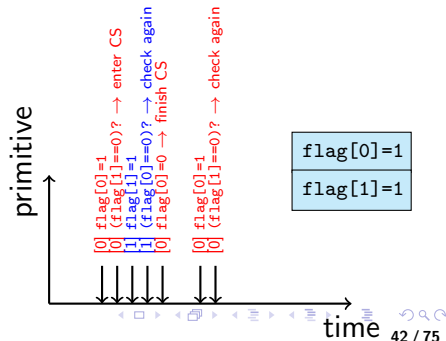
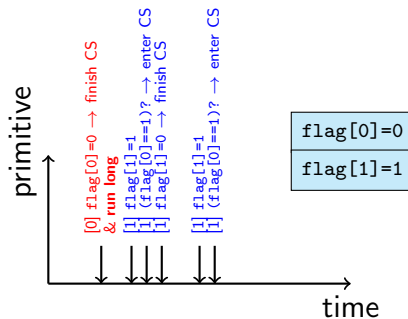
Algorithm 2 for CS

- turn cannot keep the state of readiness into CS
- → use a shared variable “flag” to keep the state of all processes
 $\text{flag}[i] = \text{true} \Rightarrow P_i \text{ is in CS}$

Process P_i

```
do {
    flag[i] = true;
    while ( flag[j] ) ;
    /* critical section code is here */
    flag[i] = false;
    /* remainder section code is here */
} while(true);
```

mutual exclusion, but not bounded waiting, and not progress



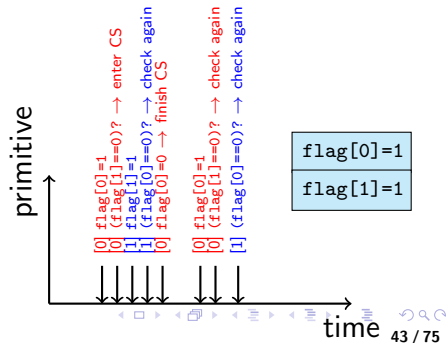

Algorithm 2 for CS

- turn cannot keep the state of readiness into CS
- → use a shared variable “flag” to keep the state of all processes
 $\text{flag}[i] = \text{true} \Rightarrow P_i$ is in CS

Process P_i

```
do {  
    flag[i] = true;  
    while ( flag[j] ) ;  
    /* critical section code is here */  
    flag[i] = false;  
    /* remainder section code is here */  
} while(true);
```

mutual exclusion, but not bounded waiting, and not progress



Peterson's algorithm for CS

Satisfy mutual exclusion

- Use both kinds of variables of Algorithm 1 & 2

Process P_i

```
do {  
    flag[i] = true;  
    turn = j;  
    while ( flag[j] && turn == j ) ;  
        /* critical section code is here */  
    flag[i] = false;  
        /* remainder section code is here */  
} while(true);
```

Bakery algorithm (1)

Peterson's algorithm is **only for 2 processes**. Bakery algorithm is for **n processes**

- Before entering its critical section, process receives a ticket number. Holder of the **smallest number** enters its CS
- If processes P_i and P_j receive the same number, if $i < j$, then P_i is served first; else P_j is served first
- The numbering scheme always generates numbers in increasing order of enumeration; i.e., 1,2,3,3,3,3,4,5,...

Bakery algorithm (2)

- A process is assigned a pair (pid,ticket #) which is ordered **lexicographically**
- Shared data
 - boolean choosing[n]
 - number[n]: ticket #

Process P_i

```
do {
    choosing[i] = true;
    number[i] = max(number[0],number[1],...,number[n-1]) + 1;
    choosing[i] = false;
    for( j=0; j<n; j++ ){
        while( choosing[j] ) ;
        while( number[j] != 0 && (number[j] < number[i] ) ;
    }
    /* critical section code is here */
    number[i] = 0;
    /* remainder section code is here */
} while(true);
```


- 1 Basic concepts
- 2 Critical-section (CS) problem
- 3 Synchronization hardware**
- 4 Software tools for synchronization
 - Mutex lock
 - Semaphores
- 5 Synchronization problems

Software-based synchronization weakness

Perterson's algorithm case

Modern computer architecture

- Multiple cores for multithreaded applications
- Processor or compiler can **reorder** the execution (if checked **no data dependency**)



Software-based solutions such as Peterson's are **not guaranteed to work** on modern computer architectures (multiple processors/cores)

There are 3 forms of hardware instructions to support solving critical-section problem.

- Memory barriers
- Hardware instructions
- Atomic variables

Memory model

A model that a computer architecture guarantees to application programs on memory.

- **Strongly ordered:** where a memory modification on one processor is immediately visible to all other processors.
- **Weakly ordered:** where modifications to memory on one processor may not be immediately visible to other processors.

Memory model

A model that a computer architecture guarantees to application programs on memory.

- **Strongly ordered:** where a memory modification on one processor is immediately visible to all other processors.
- **Weakly ordered:** where modifications to memory on one processor may not be immediately visible to other processors.

Memory models **vary by processor type**. \Rightarrow OS kernel cannot depend on any assumptions regarding the visibility of modifications to memory on a shared-memory multiprocessor.

\Rightarrow Computer architectures provide **memory barriers** to force any changes in memory to be propagated **to all other processors**.

Memory barriers - example

Thread 1

```
while (!flag)
    memory_barrier();
print x;
```

Thread 1

```
x = 100;
memory_barrier();
flag = true;
```

$\text{flag} > x$

$x > \text{flag}$

The memory barrier will ensure thread 1 will output 100.

We could place a memory barrier into Peterson's algorithm between `flag[i]` and `turn = j`.

Hardware instructions

test_and_set()

Use an **uninterruptable** (**atomic**) sequence of instructions which modifies a **shared** variable lock (initially lock=false).

test_and_set()

```
boolean test_and_set( boolean *target )
{
    boolean rv = *target;
    *target = true;
    return rv;
}
```

Mutual-exclusion by test_and_set()

```
do {
    while( test_and_set( &lock ) ) ;
    /* critical section code is here */
    lock = false;
    /* remainder section code is here */
} while (true);
```

Hardware instructions

compare_and_swap()

Swaping the contents of two words (compare_and_swap())
on a shared variable (lock) (initially lock=0)

compare_and_swap()

```
boolean compare_and_swap( int *value,  
                          int expected,  
                          int new_value )  
{  
    int temp = *value;  
    if ( *value == expected )  
        *value = new_value;  
    return temp;  
}
```

Mutual-exclusion by compare_and_swap()

```
do {  
    while( compare_and_swap( &lock, 0, 1 ) != 0 ) ;  
    /* critical section code is here */  
    lock = 0;  
    /* remainder section code is here */  
} while (true);
```

Hardware instructions

compare_and_swap()

Swaping the contents of two words (compare_and_swap())
on a shared variable (lock) (initially lock=0)

compare_and_swap()

```
boolean compare_and_swap( int *value,  
                           int expected,  
                           int new_value )  
{  
    int temp = *value;  
    if ( *value == expected )  
        *value = new_value;  
    return temp;  
}
```

Mutual-exclusion by compare_and_swap()

```
do {  
    while( compare_and_swap( &lock, 0, 1 ) != 0 ) ;  
    /* critical section code is here */  
    lock = 0;  
    /* remainder section code is here */  
} while (true);
```

The solution compare_and_swap() above is **not bounded waiting**.

Hardware instructions

bounded waiting `compare_and_swap()`

Process P_i

```
while (true) {
    waiting[i] = true;
    key = 1;
    while (waiting[i] && key == 1)
        key = compare_and_swap(&lock, 0, 1);
    waiting[i] = false;

    /* critical section */

    j = (i + 1) % n;
    while ((j != i) && !waiting[j])
        j = (j + 1) % n;

    if (j == i)
        lock = 0;
    else
        waiting[j] = false;

    /* remainder section */
}
```

- waiting initialized to false, lock initialized to 0.
- On leaving CS, P_i scans in the cyclic ordering $(i+1, i+2, \dots, n-1, 0, \dots, i-1)$ (that `waiting[j] == true`), and then set `waiting[j] == true` to allow P_j enter CS.
 \Rightarrow Bounded waiting.

Atomic variables

Increment and decrement an integer value may produce race condition.

- `compare_and_swap()` is used as **basic** building blocks for constructing synchronization tools.
- **Atomic variable**: providing atomic operations (`increment()`, `decrement()`) on basic data types such as integers and booleans.

```
void increment(atomic int *v)
{
    int temp;
    do {
        temp = *v;
    } while (temp != compare_and_swap(v, temp, temp+1));
}
```

- 1 Basic concepts
- 2 Critical-section (CS) problem
- 3 Synchronization hardware
- 4 Software tools for synchronization**
 - Mutex lock
 - Semaphores
- 5 Synchronization problems

Application programmers do **not like** using hardware-synchronization mechanisms.

OS designers must provide tools to deal with critical-section with 2 functions

- `acquire()`

```
acquire() {  
    while (!available) ; /* busy wait */  
    available = false;  
}
```

- `release()`

```
release() {  
    available = true;  
}
```

Before entering CS, programmers call `acquire()`, and call `release()` when leaving CS. Both are **atomic** which can be implemented by hardware synchronization instructions.

Application programmers do **not like** using hardware-synchronization mechanisms.

OS designers must provide tools to deal with critical-section with 2 functions

- `acquire()`

```
acquire() {  
    while (!available) ; /* busy wait */  
    available = false;  
}
```

- `release()`


Busy waiting = when a process is in CS, other processes **loop continuously** in order to enter CS (**spinlock**)

Before entering CS, programmers call `acquire()`, and call `release()` when leaving CS. Both are **atomic** which can be implemented by hardware synchronization instructions.

Semaphores

These people represent **waiting threads**.
They aren't running on any CPU core.

The bouncer represents a **semaphore**.
He won't allow threads to proceed
until instructed to do so.



(source:preashing.com)

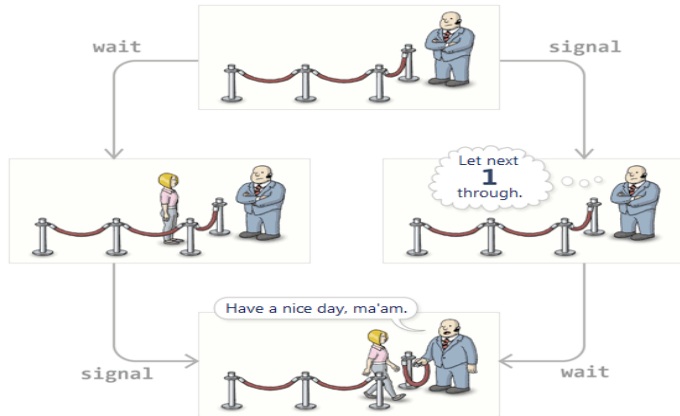
- Semaphore S: an integer variable
- can **only** be accessed by **indivisible** (**atomic**) operations

```
wait(S) {  
    while ( S <= 0 ) ; /* busy wait */  
    S --;  
}
```

```
signal(S) {  
    S ++;  
}
```

Semaphores

Calling order



(source: preshing.com)

Same outcome for different orders of calling `wait()` and `signal()`.

Semaphore usage

- Binary semaphore ($\{0,1\}$): similar to mutex locks
- Counting semaphore: to control access to resource with a **finite** number of instances
 - S initialized to number of resource instances
 - Wish to use an instance, call `wait(S)`
 - Release the resource, call `signal(S)`
 - All instances are in use, processes blocked until $S > 0$

Example of using semaphore `synch` to synchronize 2 processes (similar to `memory_barrier()`).

```
synch = 0;
```

P_1

```
S1;  
signal( synch );
```

P_2

```
wait( synch );  
S2;
```

Semaphores without busy waiting

Semaphore

Semaphore mentioned before could have **busy waiting**

Idea to avoid busy waiting

- After the checking of the semaphore fails, the process will be **put into waiting queue**
- The process is restarted by `wakeup()` when other processes send `signal()` (i.e., put the process into **ready queue**)

Semaphores without busy waiting

Implementation

Semaphore data structure

```
typedef struct {  
    int value;  
    struct process *list;  
} semaphore;
```

wait()

```
wait( semaphore *S ){  
    S->value --;  
    while ( S->value < 0 ){  
        add this process to S->list;  
        sleep();  
    }  
}
```

signal()

```
signal( semaphore *S ){  
    S->value ++;  
    if ( S->value >= 0 ) {  
        remove a process P from S->list;  
        wakeup( P );  
    }  
}
```

- block(), wakeup(): system calls

Semaphores without busy waiting

Implementation

Semaphore data structure

```
typedef struct {  
    int value;  
    struct process *list;  
} semaphore;
```

wait()

```
wait( semaphore *S ){  
    S->value --;
```

signal()

```
signal( semaphore *S ){  
    S->value ++;
```

Remarks

- Busy waiting just moved into critical section, **not completely eliminated**
- wait(), signal() must be executed **atomically** (on SMP, other techniques could be used additionally)

What is deadlock ?

P_0

```
wait(S);  
wait(Q);  
...  
signal(S);  
signal(Q);
```

P_1

```
wait(Q);  
wait(S);  
...  
signal(Q);  
signal(S);
```

- Two semaphores S and Q **initialized to 1**
- After wait(S) of P_0 and wait(Q) of P_1 executed, both processes wait for corresponding signal(s).


What is deadlock ?

P_0

```
wait(S);  
wait(Q);  
...  
signal(S);  
signal(Q);
```

P_1

```
wait(Q);  
wait(S);  
...  
signal(Q);  
signal(S);
```



- Two semaphores S and Q **initialized to 1**
- After `wait(S)` of P_0 and `wait(Q)` of P_1 executed, both processes wait for corresponding `signal()`s.

Deadlock

Two or more processes are waiting indefinitely for an event that can be caused by only one of the waiting processes.


What is deadlock ?

P_0

```
wait(S);  
wait(Q);  
...  
signal(S);  
signal(Q);
```

P_1

```
wait(Q);  
wait(S);  
...  
signal(Q);  
signal(S);
```



- Two semaphores S and Q **initialized to 1**
- After `wait(S)` of P_0 and `wait(Q)` of P_1 executed, both processes wait for corresponding `signal()`s.

Deadlock

Two or more processes are waiting indefinitely for an event that can be caused by only one of the waiting processes.

Starvation (indefinite blocking)

A process may never be removed from the semaphore queue in which it is suspended

- 1 Basic concepts
- 2 Critical-section (CS) problem
- 3 Synchronization hardware
- 4 Software tools for synchronization
 - Mutex lock
 - Semaphores
- 5 Synchronization problems**

Classic synchronization problems

- Only classic problems considered
 - Bounded-buffer problem
 - Readers-writers problem
 - Dining-philosophers problem
- Use semaphore for synchronization

Bounded-buffer problem

Semaphore data structure

```
int n;  
semaphore mutex = 1;  
semaphore empty = n;  
semaphore full = 0;
```

producer

```
do {  
    /* produce an item in      */  
    /* next produced          */  
    wait( empty );  
    wait( mutex );  
    /* add next_produced to    */  
    /* the buffer              */  
    signal( mutex );  
    signal( full );  
} while (true);
```

consumer

```
do {  
    wait( full );  
    wait( mutex );  
    /* remove an item from the */  
    /* buffer to next consumed */  
    signal( mutex );  
    signal( empty );  
    /* consume the item in     */  
    /* next_consumed           */  
} while (true);
```


Readers-writers problem

- A database shared among several processes: some are readers, some are writers
- If a writer and some other process access database **simultaneously** chaos may occur

First readers-writers problem

no reader kept waiting unless a writer has already obtained permission to use shared objects

```
semaphore rw_mutex = 1;  
semaphore mutex = 1;  
int read_count = 0;
```

Writer

```
do {  
    wait( rw_mutex );  
    .../* writing performed */...  
    signal( rw_mutex );  
} while (true);
```


Reader

```
do {  
    wait( mutex );  
    read_count ++;  
    if ( read_count == 1 )  
        wait( rw_mutex );  
    signal( mutex );  
    .../* reading performed */...  
    wait( mutex );  
    read_count --;  
    if ( read_count == 0 )  
        signal( rw_mutex );  
    signal( mutex );  
} while (true);
```

Dining-philosopher

Dining-philosopher

- Resource: 5 single chopsticks
- 5 Philosophers sitting around the table, with actions
 - Think: do nothing, not critical
 - Eat: need 2 chopsticks (left and right)



Semaphores can be used to solve dining-philosophers synchronization problem, using 5 semaphores

Dining-philosopher

Implementation

```
semaphore chopstick[5];
```

Philosopher *i*

```
do {  
    wait( chopstock[i] );  
    wait( chopstick[(i+1) % 5] );  
    /* eat eat eat */  
    signal( chopstick[i] );  
    signal( chopstick[(i+1) % 5] );  
    /* think think think */  
} while (true);
```